

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import StandardScaler
```

```
df = pd.read_csv("/content/dataset.csv (1).zip", parse_dates=[1,2])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            196441 non-null  float64
1   created_at                           197428 non-null  datetime64[ns]
2   actual_delivery_time                 197421 non-null  datetime64[ns]
3   store_id                             197428 non-null  object
4   store_primary_category               192668 non-null  object
5   order_protocol                       196433 non-null  float64
6   total_items                          197428 non-null  int64
7   subtotal                             197428 non-null  int64
8   num_distinct_items                  197428 non-null  int64
9   min_item_price                      197428 non-null  int64
10  max_item_price                       197428 non-null  int64
11  total_onshift_partners               181166 non-null  float64
12  total_busy_partners                  181166 non-null  float64
13  total_outstanding_orders             181166 non-null  float64
dtypes: datetime64[ns](2), float64(5), int64(5), object(2)
memory usage: 21.1+ MB
```

```
# feature engineering
```

```
df['time_taken'] = df['actual_delivery_time'] - df['created_at']
```

```
df['time_taken_mins'] = pd.to_timedelta(df['time_taken']) / pd.Timedelta('60s')
```

```
# some other features we can create as
```

```
df['hours'] = df['created_at'].dt.hour
```

```
df['day'] = df['created_at'].dt.dayofweek
```

```
df.head(3)
```



	market_id	created_at	actual_delivery_time	store_id	store_
0	1.0	2015-02-06 22:24:17	2015-02-06 23:27:16	df263d996281d984952c07998dc54358	
1	2.0	2015-02-10 21:49:25	2015-02-10 22:56:29	f0ade77b43923b38237db569b016ba25	
2	3.0	2015-01-22 20:39:28	2015-01-22 21:09:09	f0ade77b43923b38237db569b016ba25	

```
# drop of some useless features
```

```
df.drop(['time_taken', 'created_at', 'actual_delivery_time', 'store_id'], axis=1, inplace = 1
```

```
df.dropna(inplace = True)
```

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
Index: 176248 entries, 0 to 197427
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                            176248 non-null float64
1   store_primary_category               176248 non-null object
2   order_protocol                      176248 non-null float64
3   total_items                         176248 non-null int64
4   subtotal                           176248 non-null int64
5   num_distinct_items                 176248 non-null int64
6   min_item_price                     176248 non-null int64
7   max_item_price                     176248 non-null int64
8   total_onshift_partners              176248 non-null float64
9   total_busy_partners                 176248 non-null float64
10  total_outstanding_orders            176248 non-null float64
11  time_taken_mins                     176248 non-null float64
12  hours                              176248 non-null int32
13  day                                176248 non-null int32
dtypes: float64(6), int32(2), int64(5), object(1)
memory usage: 18.8+ MB
```

```
df['store_primary_category'].isnull().sum()
```



```
0
```

```
from sklearn.preprocessing import LabelEncoder
```

```
# Create an instance of LabelEncoder
```

```
label_encoder = LabelEncoder()
```

```
# Fit and transform the column
```

```
df['store_primary_category_encoded'] = label_encoder.fit_transform(df['store_primary_categor
```

```
# Check the result
print(df[ 'store_primary_category_encoded'].head())
```

```
⇒ 0      4
   1     46
   8     36
  14     38
  15     38
   Name: store_primary_category_encoded, dtype: int64
```

```
df.head()
```


```
⇒
```

	market_id	store_primary_category	order_protocol	total_items	subtotal	num_distir
0	1.0	american	1.0	4	3441	
1	2.0	mexican	2.0	1	1900	
8	2.0	indian	3.0	4	4771	
14	1.0	italian	1.0	1	1525	
15	1.0	italian	1.0	2	3620	

```
df.drop(['store_primary_category'], axis = 1, inplace = True)
```

```
# check some visualization
sns.heatmap(df.corr())
```


↔ <Axes: >



Inference : We observe that delivery_time does not show correlation with other feature inc

```
sns.scatterplot(x = 'time_taken_mins', y = 'subtotal', data = df)
```

↔ <Axes: xlabel='time_taken_mins', ylabel='subtotal'>




```
sns.countplot(x=df.day, palette='Set2')  
plt.show()
```


↩ <ipython-input-47-9ece2785a83d>:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

```
sns.countplot(x=df.day, palette='Set2')
```




```
sns.countplot(x=df.hours, palette='Set2')  
plt.show()
```

 <ipython-input-48-99d9c83f57f3>:1: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0.

```
sns.countplot(x=df.hours, palette='Set2')
```



```
sns.scatterplot(x = 'hours', y = 'time_taken_mins', data = df)
```

↩ <Axes: xlabel='hours', ylabel='time_taken_mins'>



df.columns

↩ Index(['market_id', 'order_protocol', 'total_items', 'subtotal',
'num_distinct_items', 'min_item_price', 'max_item_price',
'total_onshift_partners', 'total_busy_partners',
'total_outstanding_orders', 'time_taken_mins', 'hours', 'day',
'store_primary_category_encoded'],
dtype='object')

data modeling

y = df['time_taken_mins']

x = df.drop(['time_taken_mins'], axis = 1)

df.drop(['time_taken_mins'], axis = 1, inplace = True)

x_train, x_test, y_train, y_test = train_test_split(x,y,test_size = 0.2, random_state = 42)

random forest

regressor = RandomForestRegressor()

regressor.fit(x_train, y_train)



▼ RandomForestRegressor ⓘ ?

RandomForestRegressor()

```
prediction = regressor.predict(x_test)
mse = mean_squared_error(y_test, prediction)
rmse = mse **.5
print(f'mse : ', mse)
print('rmse : ', rmse)
mae = mean_absolute_error(y_test, prediction)
print('mae : ' , mae)
```



```
mse : 319.2601396303887
rmse : 17.86785212694544
mae : 11.67903887073061
```

```
r2_score(y_test, prediction)
```



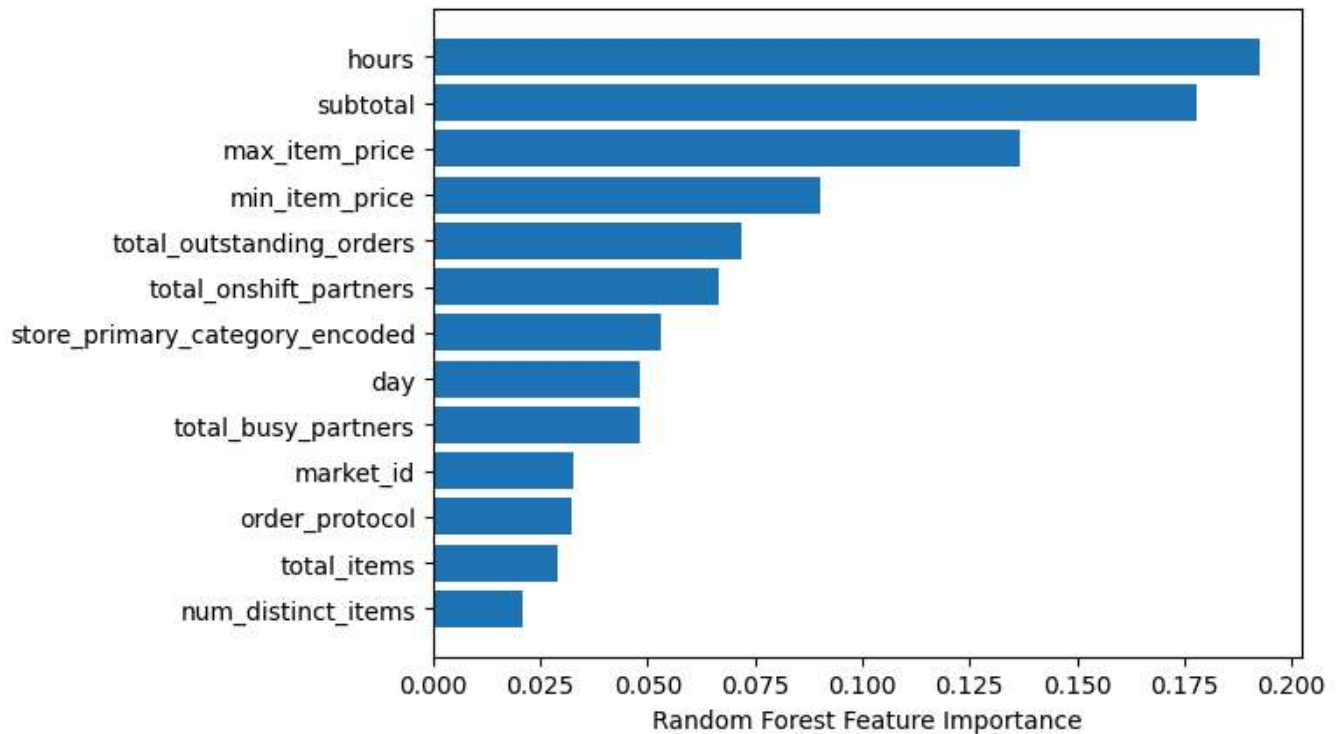
```
0.1414103997165178
```

```
def MAPE(y_actual, y_predicted):
    mape = np.mean(np.abs((y_actual - y_predicted) / y_actual)) * 100
    return mape
print ('mape : ', MAPE(y_test, prediction))
```



```
mape : 26.925932790369334
```

```
# feature importance
sorted_idx = regressor.feature_importances_.argsort()
plt.barh(df.columns[sorted_idx], regressor.feature_importances_[sorted_idx])
plt.xlabel("Random Forest Feature Importance")
plt.show()
```



```
# neural network
from sklearn import preprocessing
scaler = preprocessing.MinMaxScaler()
x_scaled = scaler.fit_transform(x)
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size = 0.2, random_state = 42)
```


```
# when to use minmax or std- max
# minmax : if range is known and there are no outliers
# std-scaler = if range ot known , then the std-scaler use also less prone to outliers
```

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
model = Sequential()
model.add(Dense(14, kernel_initializer = 'normal', activation = 'relu'))
model.add(Dense(512, activation = 'relu'))
model.add(Dense(1024, activation = 'relu'))
model.add(Dense(256, activation = 'relu'))
model.add(Dense(1, activation = 'relu'))
```


```
from tensorflow.keras.optimizers import Adam
adam = Adam(learning_rate = 0.01)
model.compile(loss = 'mse', optimizer = adam, metrics = ['mse', 'mae'])
history = model.fit(x_train, y_train, epochs = 30, batch_size = 512, verbose = 1, validation_data = (x_test, y_test))
```



Epoch 2/30

221/221  **20s** 67ms/step - loss: 1494.1469 - mae: 15.0795 - mse: 149


Epoch 3/30

221/221  **27s** 97ms/step - loss: 680.1284 - mae: 12.3244 - mse: 680


Epoch 4/30

221/221  **26s** 117ms/step - loss: 1464.8702 - mae: 12.5510 - mse: 14


Epoch 5/30

221/221  **22s** 98ms/step - loss: 1392.8450 - mae: 12.5930 - mse: 139


Epoch 6/30

221/221  **36s** 75ms/step - loss: 782.2251 - mae: 12.2248 - mse: 782


Epoch 7/30

221/221  **15s** 70ms/step - loss: 611.4301 - mae: 12.1938 - mse: 611

Epoch 8/30

221/221  **20s** 67ms/step - loss: 506.5948 - mae: 12.1799 - mse: 506

Epoch 9/30

221/221  **20s** 67ms/step - loss: 711.5001 - mae: 12.2089 - mse: 711